

Bayesian nonparametric methods for clustering

Part 2: Outcome-guided analysis

Anaïs Rouanet & Boris Hejblum

INSERM U1219, Bordeaux Population Health research Center, Bordeaux, France

CNC 2021 - April 2021, 7th

Summary of Part 1:

Nonparametric clustering

- Heterogeneous population with unknown clusters
- Cluster number estimated from the data
- Unconstrained number of clusters

Characterized by

- One or several variables (Gaussian, skew t-distributions...)

Bayesian framework

- Dirichlet process prior
- MCMC estimation

Introduction of Part 2:

Nonparametric clustering

- Heterogeneous population with unknown clusters
- Cluster number estimated from the data
- Unconstrained number of clusters


Characterized by

- One or several variables (Gaussian, skew t-distributions...)
- An outcome used as noisy surrogate for clusters (Bair, 2013)

Bayesian framework

- Dirichlet process prior
- MCMC estimation

Outcome-guided clustering



- Subgroups with specific outcome distributions and variable profiles
- Independence between profile variables and outcome conditionally on the clusters

Motivational example - Dementia study

Dementia

Syndrome (chronic or progressive nature)

Deterioration in cognitive functions

Impairment of the social and occupational functioning.



Precision medicine

Improve disease prevention and treatment

Refine patient characterisation

Predict disease risk and treatment effect at the individual level

Profile regression: methodology

Notations:

Y_i vector of individual outcome measurements

$W_i = (w_{i1}, \dots, w_{iP})^\top$ individual vector of profile variables

c_i allocation variable

Likelihood:

$$L(\theta|Y, W) = \prod_{i=1}^N \sum_g^{\infty} P(c_i = g|\theta^c) f(W_i|c_i = g; \theta_g^W) f(Y_i|c_i = g; \theta_g^Y)$$

Independence assumption between W and Y given the cluster allocations.

Profile regression: membership probability

Notations:

Y_i vector of individual outcome measurements

$W_i = (w_{i1}, \dots, w_{iP})^\top$ individual vector of profile variables

c_i allocation variable

Likelihood:

$$L(\theta|Y, W) = \prod_{i=1}^N \sum_g^{\infty} \underbrace{P(c_i = g|\theta^c)}_{\text{Membership prob.}} f(W_i|c_i = g; \theta_g^W) f(Y_i|c_i = g; \theta_g^Y)$$

Independence assumption between W and Y given the cluster allocations.

Sub-model: Dirichlet Process mixture model

$$\mathcal{A} = \sum_{g=1}^G \pi_g \delta_{\theta_g}$$

α : Concentration parameter

\mathcal{A}_0 : Base distribution

$$\pi|\alpha \sim \text{Dirichlet}(\alpha/G, \dots, \alpha/G)$$

$$c_i|\pi \sim \text{Multinom}(\pi)$$

$$\theta_{c_i}|\mathcal{A}_0 \sim \mathcal{A}_0$$

$$Y_i, W_i|c_i \sim f(\cdot; \theta_{c_i})$$

Profile regression: membership probability

Notations:

Y_i vector of individual outcome measurements

$W_i = (w_{i1}, \dots, w_{iP})^\top$ individual vector of profile variables

c_i allocation variable

Likelihood:

$$L(\theta|Y, W) = \prod_{i=1}^N \sum_g^{\infty} f(W_i|c_i = g; \theta_g^W) f(Y_i|c_i = g; \theta_g^Y)$$

Independence assumption between W and Y given the cluster allocations.

Sub-model: Dirichlet Process mixture model

$$\mathcal{A} = \sum_{g=1}^{+\infty} \pi_g \delta_{\theta_g}$$

α : Concentration parameter

\mathcal{A}_0 : Base distribution

$$\pi|\alpha \sim GEM(\alpha)$$

$$c_i|\pi \sim Multinom(\pi)$$

$$\theta_{c_i}|\mathcal{A}_0 \sim \mathcal{A}_0$$

$$Y_i, W_i|c_i \sim f(\cdot; \theta_{c_i})$$

Profile regression: profile variables

Notations:

Y_i vector of individual outcome measurements

$W_i = (w_{i1}, \dots, w_{iP})^\top$ individual vector of profile variables

c_i allocation variable

Likelihood:

$$L(\theta|Y, W) = \prod_{i=1}^N \sum_g P(c_i = g|\theta^c) \underbrace{f(W_i|c_i = g; \theta_g^W)}_{\text{Profile variables}} f(Y_i|c_i = g; \theta_g^Y)$$

Independence assumption between W and Y given the cluster allocations.

Sub-model for profile variables: given $c_i = g$,

Discrete case: $W_i^{(1)} \sim \text{Multinom}(\phi_g)$

Continuous case: $W_i^{(2)} \sim \mathcal{N}_P(\mu_g; V_g)$

Profile regression: outcome

Notations:

Y_i vector of individual outcome measurements

$W_i = (w_{i1}, \dots, w_{iP})^\top$ individual vector of profile variables

c_i allocation variable

Likelihood:

$$L(\theta|Y, W) = \prod_{i=1}^N \sum_g P(c_i = g|\theta^c) f(W_i|c_i = g; \theta_g^W) \underbrace{f(Y_i|c_i = g; \theta_g^Y)}_{\text{Outcome}}$$

Sub-model for outcome: given cluster g ,

$$Y_i \sim \mathcal{N}(\theta_g + \beta^\top X_i, \sigma_g^2)$$

Other types of outcomes:

Bernoulli, Poisson, Binomial, Categorical, Quantile, Survival

Estimated cluster-specific parameters

- MCMC Gibbs sampler: $\Theta = (\alpha, [\phi]_g, [\mu]_g, [V]_g, c, \beta)$
- Posterior similarity matrix S from pairwise co-clustering probability
- Selection of the representative clustering P^*
 - Partitioning around medoids on dissimilarity matrix $1 - S$
 - Average silhouette width

Estimated cluster-specific parameters

- MCMC Gibbs sampler: $\Theta = (\alpha, [\phi]_g, [\mu]_g, [V]_g, c, \beta)$

$$\alpha \sim f(\alpha \mid [\phi]_g, [\mu]_g, [V]_g, c, \beta)$$

- Posterior similarity matrix S from pairwise co-clustering probability
- Selection of the representative clustering P^*

Partitioning around medoids on dissimilarity matrix $1 - S$

Average silhouette width

Estimated cluster-specific parameters

- MCMC Gibbs sampler: $\Theta = (\alpha, [\phi]_g, [\mu]_g, [V]_g, c, \beta)$

$$\alpha \sim f(\alpha \mid [\phi]_g, [\mu]_g, [V]_g, c, \beta)$$

$$[\phi]_g \sim f([\phi]_g \mid \alpha, [\mu]_g, [V]_g, c, \beta) \text{ etc...}$$

- Posterior similarity matrix S from pairwise co-clustering probability
- Selection of the representative clustering P^*

Partitioning around medoids on dissimilarity matrix $1 - S$

Average silhouette width

Estimation

Cluster-specific parameters as the average over the MCMC iterations:

$$\bar{\Theta}_g^{(k)} = \frac{1}{N_g} \sum_{i|c_i^*=g} \Theta_{c_i^{(k)}}$$

$\Theta_{c_i^{(k)}}$ parameter for the cluster subject i is allocated to, at iteration k
 N_g the number of subjects in cluster g of P^* .

Estimation

Cluster-specific parameters as the average over the MCMC iterations:

$$\hat{\Theta}_g = \frac{1}{K} \sum_{k=1}^K \frac{1}{N_g} \sum_{i|c_i^*=g} \Theta_{c_i^{(k)}}$$

$\Theta_{c_i^{(k)}}$ parameter for the cluster subject i is allocated to, at iteration k
 N_g the number of subjects in cluster g of P^* .

Estimation

Cluster-specific parameters as the average over the MCMC iterations:

$$\hat{\Theta}_g = \frac{1}{K} \sum_{k=1}^K \frac{1}{N_g} \sum_{i|c_i^*=g} \Theta_{c_i^{(k)}}$$

$\Theta_{c_i^{(k)}}$ parameter for the cluster subject i is allocated to, at iteration k
 N_g the number of subjects in cluster g of P^* .

- Integration of the clustering uncertainty in the parameters estimation.
- Empirical credible intervals computed from MC sample of $\bar{\Theta}_g^{(k)}$

Convergence

Difficult to assess convergence based on parameters traces:

- Fixed effects converge quickly
- Cluster-specific parameters not traceable
- If stuck in local mode, α not good indicator (Hastie et al., 2014)

→ Marginal partition posterior $p(C|Y, W)$

- identify runs that were significantly different from others, perhaps due to convergence issues
- which run explored the higher posterior probability regions

Application to ADNI data – (Rouanet et al. 2020)

Objective : Identify subgroups of the population associated with specific cognitive evolution patterns and brain imaging profiles.

- North American Alzheimer's Disease Neuroimaging Initiative
- Sample of 199 subjects, 55 years old and over, 8-year follow-up
- **Normalised Mini-Mental State Examination** (Philipps et al., 2014)

Application to ADNI data – (Rouanet et al. 2020)

Objective : Identify subgroups of the population associated with specific cognitive evolution patterns and brain imaging profiles.

- North American Alzheimer's Disease Neuroimaging Initiative
- Sample of 199 subjects, 55 years old and over, 8-year follow-up
- **Normalised Mini-Mental State Examination** (Philipps et al., 2014)
- Profile variables:
 - **Standardised MRI brain volumes**: whole brain, ventricles, hippocampus, entorhinal cortex, fusiform gyrus and middle temporal gyrus, standardised by the intracerebral volume
 - **Gender** (1 for women, 0 for men) and **Education** (1 if greater than or equal to 16 years of education, 0 otherwise), **APOE4 status** (1 if 1 or 2 APOE4 alleles, 0 otherwise)

Model specification

Profile variables, given cluster g :

Gender, Sex, Education:

$$W_i^{(1)} \sim \text{Bernoulli}(\phi_g) \quad \text{with } \phi_g \sim \text{Dirichlet}(a_j)$$

Standardized MRI brain volumes:

$$W_i^{(2)} \sim \mathcal{N}_P(\mu_g; V_g) \quad \text{with } \mu_g \sim \mathcal{N}_P(\mu_0; V_0), \quad V_g \sim \text{InvWishart}(R_0, K_0)$$

Model specification

Profile variables, given cluster g :

Gender, Sex, Education:

$$W_i^{(1)} \sim \text{Bernoulli}(\phi_g) \quad \text{with } \phi_g \sim \text{Dirichlet}(a_j)$$

Standardized MRI brain volumes:

$$W_i^{(2)} \sim \mathcal{N}_P(\mu_g; V_g) \quad \text{with } \mu_g \sim \mathcal{N}_P(\mu_0; V_0), \quad V_g \sim \text{InvWishart}(R_0, K_0)$$

Longitudinal normalized MMSE, given cluster g :

$$Y(t_i) = f_g(t_i) + \epsilon_g(t_i) \quad \text{with } \epsilon_g \sim \mathcal{N}(0, L_{1,g})$$




$$f_g(\cdot) \sim GP(\mathbf{0}, \mathcal{K}_g(\cdot, \cdot)) \quad \text{with } \mathcal{K}_g(s, t) = L_{g,2} \exp\left(-\frac{(s-t)^2}{L_{g,3}}\right)$$

$$\exp(L_{g,l}) \sim \mathcal{N}(0, \sigma_l^2)$$


Gaussian Process prior

Given the cluster g :

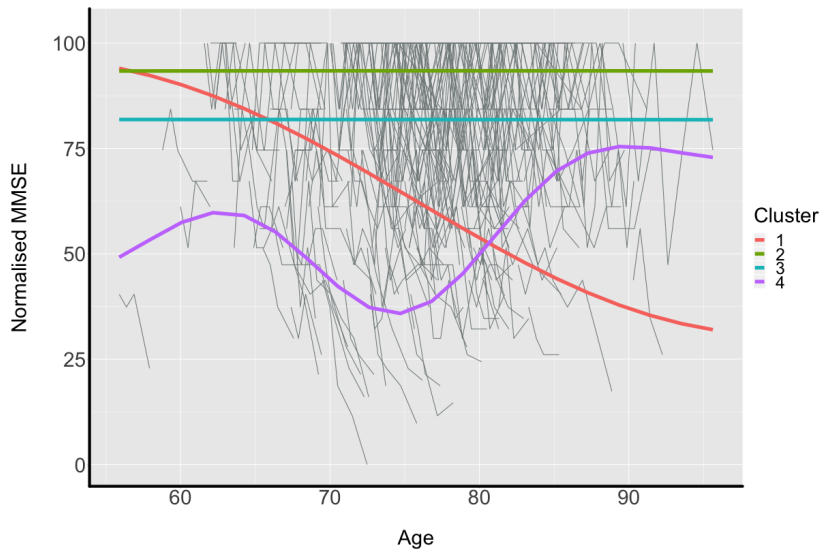
$$Y_{ij} = f_g(t_{ij}) + \epsilon_{ijg} \text{ with } \epsilon_{ijg} \sim \mathcal{N}(0, L_{g,1})$$
$$f_g(\cdot) \sim GP(m_g(\cdot), \kappa_g(\cdot, \cdot))$$



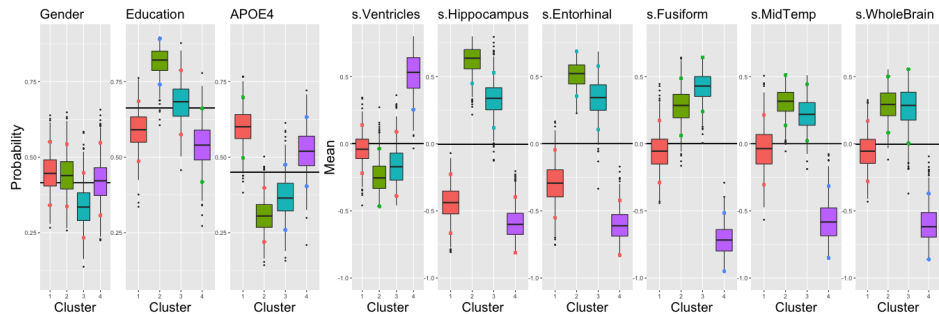
Posterior similarity matrix




Cluster-specific cognitive patterns




Cluster profiles



Cognitive pattern in cluster 4





Discussion

Conclusion

Mixture model with unconstrained number of clusters

Outcome-guided clustering analysis

Integration of different types of profile variables

Extension of the  package **PRemiuM**

Practical

Hands-on exercises on the PReMiuM package (profRegr function)

Data on cognitive decline from lcmm package

Details on DP parameters

References

Rouanet A, et al. (2020) Benefit of Bayesian clustering of longitudinal data: study of cognitive decline for precision medicine. In Bayesian Methods in Pharmaceutical Research (eds: Lesaffre E, Baio G, and Boulanger B), Chapman & Hall/CRC Biostatistics Series, ISBN: 978-1-138-74848-4 (pp. 223–242).

Rouanet A, et al. Bayesian profile regression with a longitudinal response.

Liverani S, et al. (2015). PReMiuM: An R Package for Profile Regression Mixture Models Using Dirichlet Processes. Journal of Statistical Software, 64(7),1-30.

Mueller SG, et al. (2005). Ways toward an early diagnosis in Alzheimer's disease: the Alzheimer's Disease Neuroimaging Initiative (ADNI). Alzheimers and Dementia, 1(1):55-66.

Philipps, V, et al. (2014). Normalized mini-mental state examination for assessing cognitive change in population-based brain aging studies. Neuroepidemiology, 43(1): 15–25.

Bair, E. (2013). Semi-supervised clustering methods. Wiley Interdisciplinary Reviews: Computational Statistics, 5(5): 349–361.